

# Hierarchizing Cluster Analysis: Agglomerative and Divisive Methods

Nicholas Maue, PUBPOL 542

For this analysis, we will be clustering observations based on COVID-19 deaths, as well as similar health and demographic characteristics. We will be doing both agglomerative hierarchical clustering, which determines the number of clusters from individual groups, as well as divisive clustering, which determines the number of clusters after starting from one large cluster that contains all observations. I will also compare the results of both types of clustering visually. The unit of analysis here is the county.

Reading in data

```
link='https://github.com/Public-Policy-COVID/students_merge/raw/main/Merged_data.csv'
```

```
myfile=url(link)
```

```
covid=read.csv(file=myfile)
```

Reset row names to R format

```
row.names(covid)=NULL
```

```
str(covid,width = 50,strict.width='cut')
```

```
## 'data.frame': 133 obs. of 19 variables:
## $ Number_of_beds : num 3667 0 52 553 25 ..
## $ Number_of_hospitals : num 22 0 1 6 1 1 10 1..
## $ Location : chr "Alameda_CA" "Al"..
## $ Urban_Rural_Code : chr "Large central m"..
## $ Deaths_COVID : int 573 0 31 101 12 1..
## $ Deaths_total : int 10908 0 415 2313 ..
## $ never : num 0.019 0.025 0.045..
## $ rarely : num 0.008 0.085 0.013..
## $ sometimes : num 0.055 0.088 0.099..
## $ frequently : num 0.123 0.19 0.188 ..
## $ always : num 0.795 0.612 0.655..
## $ mask_score : num 3.67 3.28 3.4 3.3..
## $ total_population : num 1671329 1129 3975..
## $ white_total_pct : num 49.3 67.9 89.7 85..
## $ black_total_pct : num 11.03 0.35 2.68 1..
## $ aian_total_pct : num 1.06 25.69 2.33 2..
## $ asian_total_pct : num 32.33 1.59 1.67 5..
```

```
## $ nhopi_total_pct      : num  0.94 0 0.29 0.29 ..
## $ multiracial_total_pct: num  5.35 4.43 3.38 4...
```

## VARIABLE PREPARATION

First, we want to include the variable for COVID deaths in the cluster analysis, but it is currently a string variable. We will need to change to numeric

```
as.numeric('Deaths_COVID')

## Warning: NAs introduced by coercion

## [1] NA
```

Now, we will choose the variables to cluster around COVID deaths. We will include total population, mask score, the number of hospital beds, the total percent of the population that is white, and the total percent of the population that is black. Although we only have total deaths and not deaths per 100K, clustering around total population should help control for variation due to population.

```
dfClus=covid[c('Number_of_beds','mask_score','Deaths_COVID','Deaths_total','Number_of_hospitals','black_total_pct','white_total_pct')]
```

```
summary(dfClus)
```

##	Number_of_beds	mask_score	Deaths_COVID	Deaths_total
##	Min. : 0.0	Min. :2.470	Min. : 0	Min. : 0
##	1st Qu.: 25.0	1st Qu.:3.301	1st Qu.: 0	1st Qu.: 0
##	Median : 131.0	Median :3.464	Median : 22	Median : 637
##	Mean : 885.4	Mean :3.428	Mean : 206	Mean : 2896
##	3rd Qu.: 553.0	3rd Qu.:3.591	3rd Qu.: 128	3rd Qu.: 2537
##	Max. :26672.0	Max. :3.822	Max. :8034	Max. :75463

##	Number_of_hospitals	black_total_pct	white_total_pct
##	Min. : 0	Min. : 0.000	Min. :49.28
##	1st Qu.: 1	1st Qu.: 0.770	1st Qu.:82.16
##	Median : 2	Median : 1.260	Median :88.64
##	Mean : 5	Mean : 2.318	Mean :85.50
##	3rd Qu.: 4	3rd Qu.: 2.620	3rd Qu.:91.84
##	Max. :112	Max. :14.770	Max. :96.13

Rescale the units into a new variable

```
dfClus=scale(dfClus)
```

```
summary(dfClus)
```

##	Number_of_beds	mask_score	Deaths_COVID	Deaths_total
##	Min. :-0.3334	Min. :-4.2726	Min. :-0.2704	Min. :-0.37704
##	1st Qu.: -0.3240	1st Qu.: -0.5659	1st Qu.: -0.2704	1st Qu.: -0.37704
##	Median : -0.2841	Median : 0.1612	Median : -0.2415	Median : -0.29411
##	Mean : 0.0000	Mean : 0.0000	Mean : 0.0000	Mean : 0.00000
##	3rd Qu.: -0.1252	3rd Qu.: 0.7277	3rd Qu.: -0.1024	3rd Qu.: -0.04674

```
## Max. : 9.7118 Max. : 1.7581 Max. :10.2736 Max. : 9.44771
## Number_of_hospitals black_total_pct white_total_pct
## Min. :-0.44686 Min. :-0.8976 Min. :-3.8920
## 1st Qu.:-0.35749 1st Qu.:-0.5994 1st Qu.:-0.3585
## Median :-0.26812 Median :-0.4097 Median : 0.3379
## Mean : 0.00000 Mean : 0.0000 Mean : 0.0000
## 3rd Qu.:-0.08937 3rd Qu.: 0.1169 3rd Qu.: 0.6818
## Max. : 9.56284 Max. : 4.8214 Max. : 1.1428
```

We will set Location as the row names, which will allow us to look at cluster results for each county.

```
row.names(dfClus)=covid$Location
head(dfClus)
```

```
##      Number_of_beds mask_score Deaths_COVID Deaths_total
## Alameda_CA      1.0476322  1.0666781    0.4816583    1.04310240
## Alpine_CA      -0.3334445 -0.6640201   -0.2703586   -0.37704284
## Amador_CA      -0.3138601 -0.1465949   -0.2296736   -0.32301275
## Butte_CA       -0.1251719 -0.2090427   -0.1378042   -0.07590643
## Calaveras_CA   -0.3240289 -0.6104934   -0.2546096   -0.32691854
## Colusa_CA      -0.3153666  0.1790262   -0.2546096   -0.36194046
##      Number_of_hospitals black_total_pct white_total_pct
## Alameda_CA      1.51932965    3.3732540   -3.89196747
## Alpine_CA      -0.44686166   -0.7620594   -1.88666366
## Amador_CA      -0.35748933    0.1401204    0.44640953
## Butte_CA       0.08937233   -0.1618969    0.01654805
## Calaveras_CA   -0.35748933   -0.4794022    0.58611451
## Colusa_CA      -0.35748933   -0.3903458    0.60223432
```

```
set.seed(999) ##This is for replicability of results
```

Determine the distance method and compute distance matrix

```
library(cluster)
dfClus_D=cluster::daisy(x=dfClus)
```

## HIERARCHIZING AGGLOMERATIVE

Set the number of clusters

```
NumCluster=4
```

Next, apply the function:

```
library(factoextra)

## Loading required package: ggplot2

## Welcome! Want to learn more? See two factoextra-related books at https://g
oo.gl/ve3WBa
```

```
res.agnes= hcut(dfClus_D,
                k = NumCluster,isdiss=T,
                hc_func='agnes',
                hc_method = "ward.D2")
```

Cluster:

```
covid$agn=as.factor(res.agnes$cluster)
```

Let's check the first cluster results

```
covid[covid$agn==1, 'Location']

## [1] "Alameda_CA"      "Contra Costa_CA"  "Orange_CA"
## [4] "Riverside_CA"     "Sacramento_CA"   "San Bernardino_CA"
## [7] "San Diego_CA"     "San Francisco_CA" "San Mateo_CA"
## [10] "Santa Clara_CA"  "Solano_CA"       "King_WA"
```

Let's check the results through a table

```
table(covid$agn)

##
##  1  2  3  4
## 12 45 75  1
```

The results indicate that Cluster 4 has only one observation.

```
covid[covid$agn==4, 'Location']

## [1] "Los Angeles_CA"
```

Los Angeles County appears to be the sole observation.

King County's cluster:

```
covid[covid$Location=="King_WA", 'agn']

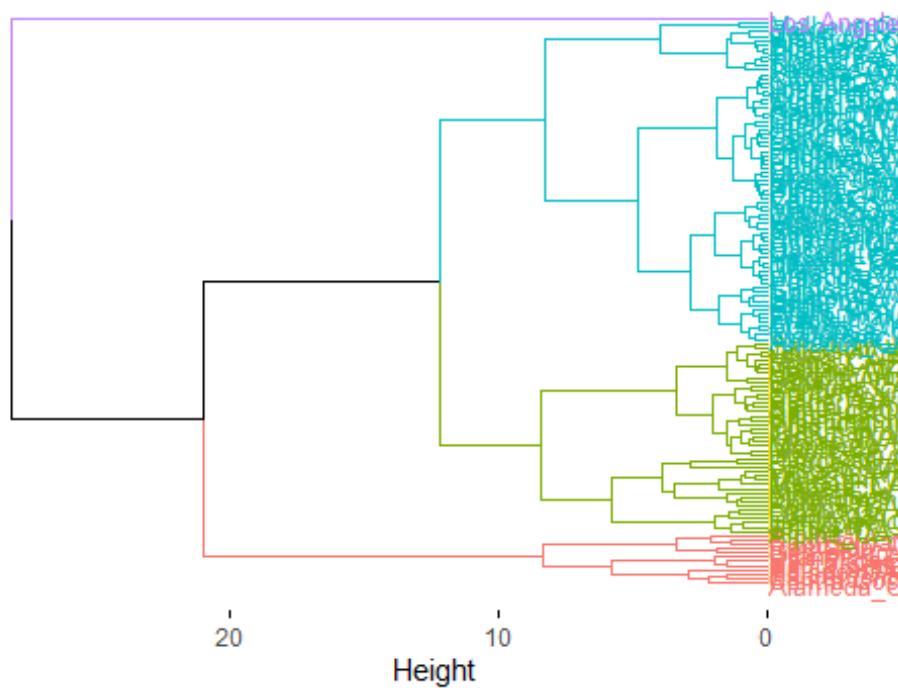
## [1] 1
## Levels: 1 2 3 4
```

## VISUALIZING AGGLOMERATIVE RESULTS

We will produce a dendrogram of the cluster results

```
fviz_dend(res.agnes,k=NumCluster, cex = 0.7, horiz = T)
```

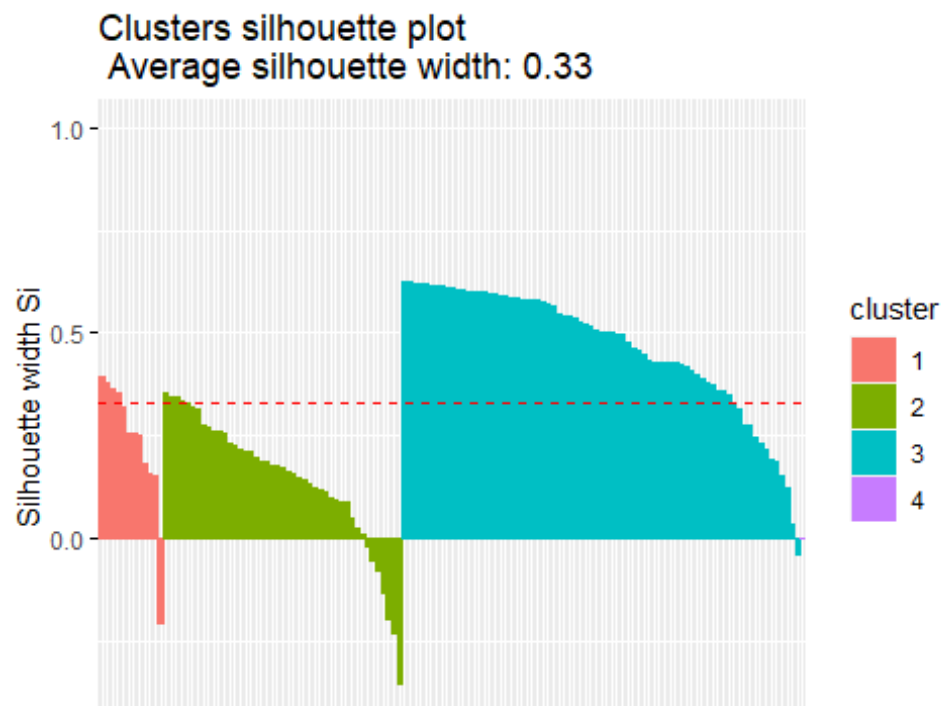
## Cluster Dendrogram



Visualize results with a silhouette plot

```
library(factoextra)
fviz_silhouette(res.agnes)
```

```
##   cluster size ave.sil.width
## 1      1    12         0.24
## 2      2    45         0.14
## 3      3    75         0.47
## 4      4     1         0.00
```



It appears that clusters 1,2, and 3 each have negative silhouettes, which means these are poorly clustered

Saving individual silhouettes:

```
agnEval=data.frame(res.agnes$silinfo$widths)
head(agnEval)
```

	cluster	neighbor	sil_width
## Alameda_CA	1	2	0.3919787
## King_WA	1	2	0.3804824
## Sacramento_CA	1	2	0.3653891
## San Bernardino_CA	1	2	0.3530916
## San Diego_CA	1	2	0.3198629
## Riverside_CA	1	2	0.2560581

Check the observations with negative silhouettes

```
agnEval[agnEval$sil_width<0,]
```

	cluster	neighbor	sil_width
## San Mateo_CA	1	2	-0.20528803
## Lassen_CA	2	3	-0.01850048
## San Joaquin_CA	2	1	-0.05384687
## Jefferson_WA	2	3	-0.07849399
## Del Norte_CA	2	3	-0.13087586
## San Juan_WA	2	3	-0.19746628
## Ferry_WA	2	3	-0.23240216

```
## Okanogan_WA      2      3 -0.35433693
## San Benito_CA    3      2 -0.04121742
```

In total, there are nine observations that are poorly clustered: One in cluster 1, Seven in Cluster 2, and One in Cluster 3

## HIERARCHIZING DIVISIVE METHOD

Apply the function (we will use the same number of clusters, 4, as the agglomerative method)

```
library(factoextra)

res.diana= hcut(dfClus_D, k = NumCluster,
               hc_func='diana',
               hc_method = "ward.D")
```

Clustering

```
covid$dia=as.factor(res.diana$cluster)
```

Querying the data frame

```
covid[covid$dia==1, 'Location']

## [1] "Alameda_CA"      "Contra Costa_CA" "Fresno_CA"      "Sacramento_
## [5] "San Francisco_CA" "San Joaquin_CA"   "San Mateo_CA"    "Santa Clara
## [9] "Solano_CA"       "Pierce_WA"
```

Check the results by each cluster:

```
table(covid$dia)

##
##  1   2   3   4
## 10 117   1   5
```

The results indicate that Cluster 2 has the large majority of observations. Let's check that one:

```
covid[covid$dia==2, 'Location']

## [1] "Alpine_CA"      "Amador_CA"      "Butte_CA"
## [4] "Calaveras_CA"   "Colusa_CA"      "Del Norte_CA"
## [7] "El Dorado_CA"   "Glenn_CA"       "Humboldt_CA"
## [10] "Imperial_CA"    "Inyo_CA"        "Kern_CA"
## [13] "Kings_CA"       "Lake_CA"        "Lassen_CA"
## [16] "Madera_CA"      "Marin_CA"       "Mariposa_CA"
## [19] "Mendocino_CA"   "Merced_CA"      "Modoc_CA"
## [22] "Mono_CA"        "Monterey_CA"    "Napa_CA"
## [25] "Nevada_CA"      "Placer_CA"      "Plumas_CA"
```

```
## [28] "San Benito_CA"      "San Luis Obispo_CA" "Santa Barbara_CA"
## [31] "Santa Cruz_CA"      "Shasta_CA"          "Sierra_CA"
## [34] "Siskiyou_CA"        "Sonoma_CA"          "Stanislaus_CA"
## [37] "Sutter_CA"          "Tehama_CA"          "Trinity_CA"
## [40] "Tulare_CA"          "Tuolumne_CA"        "Ventura_CA"
## [43] "Yolo_CA"            "Yuba_CA"            "Baker_OR"
## [46] "Benton_OR"          "Clackamas_OR"       "Clatsop_OR"
## [49] "Columbia_OR"        "Coos_OR"            "Crook_OR"
## [52] "Curry_OR"          "Deschutes_OR"       "Douglas_OR"
## [55] "Gilliam_OR"         "Grant_OR"           "Harney_OR"
## [58] "Hood River_OR"      "Jackson_OR"         "Jefferson_OR"
## [61] "Josephine_OR"       "Klamath_OR"         "Lake_OR"
## [64] "Lane_OR"            "Lincoln_OR"         "Linn_OR"
## [67] "Malheur_OR"         "Marion_OR"          "Morrow_OR"
## [70] "Multnomah_OR"       "Polk_OR"            "Sherman_OR"
## [73] "Tillamook_OR"       "Umatilla_OR"        "Union_OR"
## [76] "Wallowa_OR"         "Wasco_OR"           "Washington_OR"
## [79] "Wheeler_OR"         "Yamhill_OR"         "Adams_WA"
## [82] "Asotin_WA"          "Benton_WA"          "Chelan_WA"
## [85] "Clallam_WA"         "Clark_WA"           "Columbia_WA"
## [88] "Cowlitz_WA"         "Douglas_WA"         "Ferry_WA"
## [91] "Franklin_WA"        "Garfield_WA"        "Grant_WA"
## [94] "Grays Harbor_WA"    "Island_WA"          "Jefferson_WA"
## [97] "Kitsap_WA"          "Kittitas_WA"        "Klickitat_WA"
## [100] "Lewis_WA"           "Lincoln_WA"         "Mason_WA"
## [103] "Okanogan_WA"        "Pacific_WA"         "Pend Oreille_WA"
## [106] "San Juan_WA"        "Skagit_WA"          "Skamania_WA"
## [109] "Snohomish_WA"       "Spokane_WA"         "Stevens_WA"
## [112] "Thurston_WA"        "Wahkiakum_WA"       "Walla Walla_WA"
## [115] "Whatcom_WA"         "Whitman_WA"         "Yakima_WA"
```

Let's check King County:

```
covid[covid$Location=="King_WA" , 'dia']
```


```
## [1] 4
## Levels: 1 2 3 4
```

Produce silhouettes to visualize results and report average silhouettes

```
library(factoextra)
fviz_silhouette(res.diana)

## cluster size ave.sil.width
## 1      1   10      0.25
## 2      2  117      0.63
## 3      3    1      0.00
## 4      4    5      0.43
```





Cluster 2 has a negative silhouette, meaning it is poorly clustered.

Next we will save silhouettes

```
diaEval=data.frame(res.diana$silinfo$widths)
head(diaEval)
```

	cluster	neighbor	sil_width
## San Joaquin_CA	1	4	0.4104780
## Contra Costa_CA	1	4	0.4078738
## Solano_CA	1	4	0.3537789
## San Francisco_CA	1	4	0.3027197
## Alameda_CA	1	4	0.2676310
## Sacramento_CA	1	4	0.2209684

Let's check the poorly clustered silhouette in Cluster 2

```
diaEval[diaEval$sil_width<0,]
```

	cluster	neighbor	sil_width
## Multnomah_OR	2	1	-0.04078789

It looks like Multnomah County is the poorly clustered result

COMPARING AGGLOMERATIVE AND DIVISIVE CLUSTERS

Prepre a bidimensional map


```
projectedData = cmdscale(dfClus_D, k=2)
```

Save coordinates to original data frame

```
covid$dim1 = projectedData[,1]
covid$dim2 = projectedData[,2]
```

Map the Clusters

```
base= ggplot(data=covid,
             aes(x=dim1, y=dim2,
                 label=Location))
base + geom_text(size=2)
```



Plot the Agglomerative Results


```
agnPlot=base + labs(title = "AGNES") + geom_point(size=2,
                                                    aes(color=agn),
                                                    show.legend = F)
```

Plot the Divisive Results

```
diaPlot=base + labs(title = "DIANA") + geom_point(size=2,
                                                    aes(color=dia),
                                                    show.legend = F)
```

Let's look at the visual results

```
library(ggpubr)
ggarrange(agnPlot, diaPlot, ncol = 3)
```




From the visual plots, the results for both hierarchical clustering methods appear to be pretty consistent.

We can label the two hierarchical clustering plots


```
library(ggrepel)
agnPlot + geom_text_repel(size=3,aes(label=Location))

## Warning: ggrepel: 120 unlabeled data points (too many overlaps). Consider
## increasing max.overlaps
```



```
diaPlot + geom_text_repel(size=3,aes(label=Location))
```

```
## Warning: ggrepel: 120 unlabeled data points (too many overlaps). Consider  
## increasing max.overlaps
```




It looks like there are too many overlaps for the large majority of names to appear. Instead, we will need to try and label anomalies from the agn and dia plots

```
LABEL=ifelse(diaEval$sil_width<0, covid$Location,"")
```

```
diaPlot + geom_text_repel(aes(label=LABEL))
```


```
## Warning: ggrepel: 1 unlabeled data points (too many overlaps). Consider
## increasing max.overlaps
```



```
LABEL=ifelse(agnEval$sil_width<0, covid$Location,"")
```

```
agnPlot + geom_text_repel(aes(label=LABEL))
```

```
## Warning: ggrepel: 8 unlabeled data points (too many overlaps). Consider  
## increasing max.overlaps
```



It looks like we still get some overlaps